

Borla

Project Documentation

Student: **George Boadu Junior**

Student ID: **22427354**

Course: CSCD 602 — Advanced Software Engineering, University of Ghana

Document: Project_Documentation.pdf · Version 1.0 · 13 August 2026

1. Project title

Borla — a real-time matchmaker connecting Ghanaian households with nearby roaming waste collectors. ("Borla" is Ghanaian slang for rubbish/waste.)

2. Problem statement

In many Ghanaian neighbourhoods, formal municipal waste collection is unreliable or absent, so households depend on informal, roaming waste collectors (tricycle/pickup operators) who work a patch of streets looking for waste to collect for a fee. There is no shared channel connecting the two sides: a household with waste ready has no way to signal it beyond hoping a collector happens to pass, and a collector has no way to see where waste is actually waiting without driving around speculatively. Both sides waste time and, for collectors, fuel/effort — a classic two-sided discovery problem that ride-hailing-style matchmaking solves in other domains.

3. Aim and objectives

Aim: build and deploy a functional, end-to-end matchmaking platform that lets a household announce or directly request a waste pickup, and lets a collector discover and respond to that demand in real time — while demonstrating disciplined requirements engineering, effort estimation, design, implementation, testing, technical-debt management, and deployment practice within a hard time constraint.

Objectives:

1. Model the matching problem as two independent mechanisms — a stateless broadcast and a stateful direct request — rather than one over-complicated dispatcher.
2. Build a working two-sided trust layer (ratings) so both parties have a reason to behave well.
3. Give operations staff (admin) real levers: verify collectors, moderate content, watch live density, and retune behaviour without a redeploy.
4. Apply an explicit, justified effort-estimation technique to *scope* the build, not just to describe it after the fact.
5. Ship something genuinely running at a public URL, not a mock-up — with real, executable automated tests and a maintained technical-debt register.

4. Stakeholders

Stakeholder	Interest
Households	Fast, low-effort way to get waste collected; trust that a collector who says they're coming actually will
Collectors	Reliable visibility into nearby demand; protection from wasted trips to stale/cleared pins; a fair way to build a reputation
Platform operators (Admin)	Ability to onboard/verify collectors, keep content safe, see whether the marketplace is healthy per neighbourhood
Course examiner	Evidence of systematic Advanced Software Engineering practice across the full lifecycle

5. Requirements analysis

Requirements were gathered by extending an existing, larger architectural vision (`borla-technical-design.md` , authored ahead of this build as the target end-state) down to what is actually implementable, deployable, and testable inside this exam's constraints. The extension/reduction process:

1. Took every functional area of the original design (two matching planes, presence, trust layer, admin, AI) as the candidate requirement pool.
2. Removed anything that depends on infrastructure or accounts that cannot be provisioned inside the exam (native Android background location, a funded SMS gateway — later *partially* resolved once real GiantSMS credentials became available mid-build, see `Technical_Debt_Plan.md` TD-02 — managed Redis, a second admin application).
3. Classified everything that remained with MoSCoW (`SRS.md` §6).
4. Ran Use Case Points estimation (§7 below) against the *remaining* scope to confirm the prioritisation was still necessary even after the initial cut — it was, by roughly 16–80× the nominal time-box, which is exactly the signal the exam brief anticipates.

Full functional/non-functional requirement tables are in `SRS.md` §4–§5.

6. Software Requirements Specification (summary)

See `SRS.md` for the complete document. In summary: three human actors (household, collector, admin) plus two system actors (the Gemini moderation API, the cron scheduler); 26 functional requirements spanning auth, the broadcast plane, the request plane, the trust layer, and the admin console; ten non-functional requirements covering security, privacy, reliability, usability, and fail-safe moderation.

7. Software effort estimation (summary)

Technique: Use Case Points, cross-checked with expert top-down estimation. Full working in `SRS.md` §7.

Headline numbers:

Method	Result	Interpretation
UCP (216 UUCP × 1.05 TCF × 0.845 EF × 20h/UCP)	≈ 3,833 person-hours (≈ 22.6 person-months)	Cost to build this exact scope to full commercial rigour
Expert top-down (module-by-module + 15% overhead)	≈ 770 person-hours (≈ 4.5 person-months solo)	Cost for one experienced engineer to build a lean-but-real MVP
Nominal exam window	48 hours	Neither estimate is remotely close to this — by design (see SRS §7.7)

How this shaped scope: since both estimates dwarf the available time regardless of how aggressively the *original* design was already cut, the estimate's job was to decide *how* to spend the tiny fraction of either number actually available: breadth (a complete, working, thin slice of every Must-Have mechanism) rather than depth (fully hardening a narrower subset). The resulting shortfall — against either estimate — is exactly what §12 (Technical debt) below catalogues.

8. System analysis

Actors and their goals: a household wants waste gone with minimum effort and no wasted trust; a collector wants reliable, current demand signals and protection from stale pins; an admin wants levers to keep the marketplace healthy and safe without needing a redeploy for every tuning change.


Core domain insight carried over from the original design: the two matching mechanisms have opposite integrity properties on purpose. A **broadcast** is stateless and lossy — no collector "owns" a pin, so there is no claim race to resolve, at the cost of possible wasted trips (mitigated by a strong clear-prompt and short TTL). A **request** is stateful and narrow — exactly one collector is targeted, so the only correctness hazard is a stale accept/reject arriving after the request already resolved, which is closed with a conditional `WHERE status IN (...)` update rather than a lock. Keeping these two mechanisms structurally separate (different tables, different state machines) is why neither had to inherit the other's complexity.

Key non-functional pressure: every mutation that touches trust (reviews, contact reveal, admin actions) needed an explicit integrity rule *enforced in the schema*, not just in application code — see the `UNIQUE(author_id, request_id)` / `UNIQUE(author_id, broadcast_id)` constraints and the `CHECK` that a review references exactly one interaction type.

9. System design

9.1 Architecture

One monorepo, one Postgres database, one Render Web Service serving both the built React app and the `/api` + Socket.IO endpoints from the same origin (no CORS, no second service to keep alive on a free tier):




Deliberate substitutions from the original design (each is a named technical-debt item, not an oversight): PostGIS/Postgres directly instead of a Redis hot-path + Postgres cold-path split (TD-05); `node-cron` sweeps instead of BullMQ+Redis jobs (TD-05); Leaflet+OSM raster tiles instead of MapLibre+vector tiles (TD-08); one Express

service instead of separate household/collector/admin apps (TD-07); phone+password admin login instead of email+2FA (TD-06).


9.2 Data design (ER overview)

The full schema is `server/migrations/001_init.sql` (12 tables). Core entities and their relationships:




Supporting tables not pictured: `otp_codes` (OTP hashes/expiry), `broadcast_notifications` (who was notified, for "Did they come?" attribution), `pickup_confirmations`, `review_replies` (≤1 per review), `moderation_flags` (AI/user-report queue), `audit_log` (every admin action), `app_config` (live-tunable settings).


9.3 Use-case diagram



9.4 Sequence: broadcast fan-out (design §3.1/§7)




9.5 State machine: request plane (design §3.2)



Guard on every transition out of REQUESTED/SEEN:

```
UPDATE requests SET status = '<new>' WHERE id = $1 AND status IN ('requested','seen')
0 rows affected => already resolved => respond 409, never resurrect a stale transition.
```

9.6 Activity: two-sided review, double-blind release (design §16)



10. Implementation

10.1 Technology stack

Layer	Choice	Why
Server language/framework	Node.js + Express + TypeScript	Same language across stack; Express is small enough to reason about every middleware
Database	PostgreSQL 16 + PostGIS	Real geospatial queries (<code>ST_DWithin</code> , <code>ST_Distance</code>) without hand-rolled Haversine math

Layer	Choice	Why
Realtime	Socket.IO	JWT-authable, room-per-user, graceful fallback to polling already built into the protocol
Scheduling	<code>node-cron</code>	No extra infrastructure (Redis) for a single-instance deploy
Frontend	React + Vite + TypeScript	Fast dev loop; one codebase serves all three roles via role-routing
Maps	Leaflet + OpenStreetMap raster tiles	Zero-config, no tile-provider account
AI	Google Gemini API (optional)	Free tier for moderation; degrades gracefully without a key
SMS	GiantSMS (optional)	Real OTP delivery when configured; falls back safely (see Technical_Debt_Plan TD-02)
Auth	JWT (access+refresh) + bcrypt	Stateless API auth; OTP/password hashes never stored in plaintext
Deployment	Render (one Web Service + one Postgres)	Single free-tier footprint, <code>render.yaml</code> blueprint

10.2 What was actually built

Every functional requirement tagged Must-Have or Should-Have in `SRS.md` §6 is implemented and exercised by the automated test suite: phone/OTP auth with real SMS delivery attempt, admin password login, presence toggle/heartbeat with self-healing offline detection, the full broadcast plane (create/fan-out/clear/auto-expire), the full request plane (create/seen/accept/reject/auto-timeout with idempotent transitions), reveal-on-accept masked contact, two-sided reviews with real-interaction enforcement, AI-or-manual moderation, one capped public reply, double-blind release with rolling rating aggregates, and the full admin console (verify/suspend/reinstate, moderation queue, audit log, live stats, live ops map, live config tuning).

10.3 Code organisation

```

server/src/modules/{auth,presence,broadcasts,requests,reviews,admin}/routes.ts  - one router per domain
server/src/jobs/index.ts                                                       - all 4 cron sweeps
server/src/ai/moderation.ts, server/src/utils/{otp,sms,jwt,quietHours}.ts      - isolated, unit-testable logic
client/src/pages/{Login,HouseholdHome,CollectorHome,AdminDashboard,Profile}.tsx - one screen per role/concern
client/src/components/{MapView,ReviewForm,ProtectedRoute}.tsx                 - shared, reusable

```

10.4 Security controls actually implemented

JWT-gated routes with role middleware; bcrypt-hashed OTP codes and admin passwords; parameterised SQL everywhere (no injection surface); per-IP rate limiting on all `/api/` traffic plus a dedicated OTP-request cap; ownership checks on every mutating route (a household can only clear *its own* broadcast, etc.); phone numbers stripped from API responses until a request is accepted; fail-closed moderation (no verdict $\Rightarrow$ stays hidden).

11. Testing (summary)

38/38 automated tests passing (34 server — unit + Supertest integration against a real PostgreSQL+PostGIS instance; 4 client — React Testing Library) at time of submission, plus a scripted manual system/UAT pass and a security/usability review. Two real defects were caught and fixed during development (a broken first-time-signup

code path, and a review-reply status check gap) — full detail, every test case, and both defect write-ups are in `Testing_Report.md`.

12. Technical debt

Thirteen tracked items (`Technical_Debt_Plan.md`), each with Debt→Cause→Impact→Priority→ Resolution. One is ● Critical (admin has no 2FA), six are ● Scheduled — including the GiantSMS OTP integration, which was **confirmed live in production** (the gateway accepted a real send request end-to-end) but not yet confirmed to a real handset — and the rest are ● Acceptable deliberate trade-offs. The single largest is the missing native background-location collector app (TD-03) — the original design's own #1 risk — deliberately left as the biggest future-evolution item rather than attempted unsafely inside the exam window. Full register, priorities, and a phased repayment plan are in `Technical_Debt_Plan.md`.

13. Deployment

Deployed to **Render**: one free-tier Web Service (`buildCommand: npm ci --include=dev && npm run build` , `startCommand: npm run start -w server`) serving both the compiled API and the built React client from one origin, plus one free-tier managed Postgres with the PostGIS extension enabled by the first migration. `render.yaml` in the repository root fully describes the topology (Infrastructure-as-Code, not click-ops) including which secrets (`GEMINI_API_KEY` , `GIANTSMS_API_TOKEN` , `GIANTSMS_SENDER_ID`) must be set manually in the dashboard (`sync: false`) versus generated automatically (`JWT_*_SECRET` , `generateValue: true`) or wired from the managed database (`DATABASE_URL` , `fromDatabase`). See `Deployment_and_Source_Links.txt` for the live URL and credentials, and `User_Manual.md` §6 for exact redeploy/rollback steps.

14. User manual (summary)

Full walkthroughs for all three roles — including how to read the on-screen OTP fallback if SMS delivery isn't available — are in `User_Manual.md`. In short: pick a role at first login, a household broadcasts with one tap or searches nearby collectors to request directly, a collector toggles online and responds to what appears; an admin logs in with a phone+password at the same `/login` screen (toggle to "Admin") and lands on `/admin`.

15. Maintenance strategy

Type	Approach
Corrective	GitHub issue → reproduce with an integration test that fails → fix → test passes → deploy. The existing 34 server tests are the regression net.
Adaptive	Config changes (radius, TTLs, timeouts) go through <code>app_config</code> and the admin UI — no redeploy needed for the most likely "the environment changed" adjustments.
Perfective	Tracked as the technical-debt repayment plan (<code>Technical_Debt_Plan.md</code> §4) — i18n, offline shell, deeper test coverage.
Preventive	<code>npm audit</code> run before any dependency bump; the health-check endpoint (<code>/api/health</code>) lets Render auto-restart a wedged instance.
Security updates	Dependencies pinned with caret ranges; <code>npm audit</code> reviewed monthly if the project continues; JWT secrets rotated by regenerating the Render env vars (stateless tokens, no migration needed).

Type	Approach
Dependency updates	Renovate/Dependabot recommended once this leaves the exam context — not configured now, itself a minor debt item.
Performance	PostGIS GIST index already in place for the two hot geo-queries; would add a read replica for the admin analytics reads before scaling further (per the original design's own §17 guidance).
Scalability	The single-instance ceiling (TD-05) is the first thing to remove before real growth — Redis + BullMQ + Socket.IO Redis adapter, in that order.
New features / user feedback	Photo→waste-type classification (TD-11) and native background location (TD-03) are the two most-requested-shaped gaps based on the original design's own accessibility and reliability goals.
Technology changes	Node/Postgres/React are all mainstream, actively maintained — no forced-migration risk foreseen in the near term.

16. Future evolution

Phased roughly as the original design's own roadmap (`borla-technical-design.md` §15), adjusted for what this build already covers:

- **v1.1:** real SMS delivery confirmed end-to-end (finish TD-02), admin 2FA (TD-06).
- **v1.2:** Gemini key provisioned in production, CI pipeline on every push (TD-01, TD-13).
- **v1.3:** native Capacitor collector app with background geolocation (TD-03) — tested on real low-end Android hardware, per the original design's own risk mitigation.
- **v2.0:** Redis + BullMQ + Socket.IO Redis adapter for horizontal scale (TD-05), i18n in Twi/Ga (TD-09), offline PWA shell (TD-04), photo→waste-type classification (TD-11).
- **Beyond:** the original design's own explicitly-deferred set — in-app payments/escrow, live turn-by-turn tracking, auto-fallback dispatch, provider call-masking, and expansion beyond a single pilot neighbourhood.

17. Limitations

This is a web-only build: collectors only report position while their browser tab is open and foregrounded (no native background service — TD-03), which the original design itself identifies as the single biggest reliability risk for a real deployment. OTP delivery via GiantSMS is confirmed live in production at the HTTP layer but not yet confirmed to a real handset (TD-02). Admin authentication has no 2FA (TD-06). Test coverage is broad across every module but not exhaustive within any one of them (TD-10). The system has only ever run against a handful of seeded/manually-created accounts — it has not been load-tested.

18. Conclusion

Borla demonstrates that a genuinely two-sided, stateful/stateless-hybrid matching platform with a real trust layer and admin console can be specified, estimated, designed, built, tested, documented, and deployed inside a hard time-box — provided the estimation step is used honestly to *cut scope*, not to justify skipping it. The most valuable engineering decision in this build was not any single line of code but the early recognition (via UCP) that the original design was 16–80× too large for the window available, which forced a breadth-first, thin-vertical-slice strategy instead of a narrower, more polished one. What that strategy left unfinished is not hidden — it is named, prioritised, and scheduled in `Technical_Debt_Plan.md`.

19. References

- Course specification: *CSCD 602 Advanced Software Engineering Project Exams.pdf*, University of Ghana, Department of Computer Science (Examiner: Prof. Solomon Mensah).
- Original architectural design: `borla-technical-design.md` (authored ahead of this build).
- UI design system: `borla_UI_design.html` (colour tokens, typography, and component patterns reused directly in `client/src/styles/`).
- Karner, G. (1993). *Resource Estimation for Objectory Projects* — Use Case Points method.
- Express — <https://expressjs.com> · PostgreSQL/PostGIS — <https://postgis.net> · Socket.IO — <https://socket.io> · React/Vite — <https://vitejs.dev> · Leaflet/OpenStreetMap — <https://leafletjs.com>, <https://www.openstreetmap.org> · node-cron — <https://www.npmjs.com/package/node-cron> · Google Gemini API — <https://ai.google.dev> · GiantSMS — <https://giantsms.com/developer> · Render — <https://render.com/docs>.